Spark doesn't need an introduction, but just in case, it extends Hadoop's distributed / parallel computing paradigm by using both live memory in addition to disk storage. This makes for a much faster tool. As of version 1.4, **SparkR** is included in Apache's Spark build. We'll be using version 1.5 in this walk-through. Here we'll create our **AWS EC2** instance.

- AWS Instance
- VPC
- EC2
- Connecting to EC2

Cluster Launcher

AWS Instance

In order to approach this from the same vantage point, we'll use a small EC2 instance to launch our Spark clusters. You will need an amazon **AWS** account and the ability to **Secure Shell Tunneling (SSH)** into AWS (more on this later). Keep in mind that AWS instances cost money and the more clusters you need, the more money Amazon will charge. So, always remember to "Stop" or "Terminate" you instances when not needed! First, sign into the AWS Console:

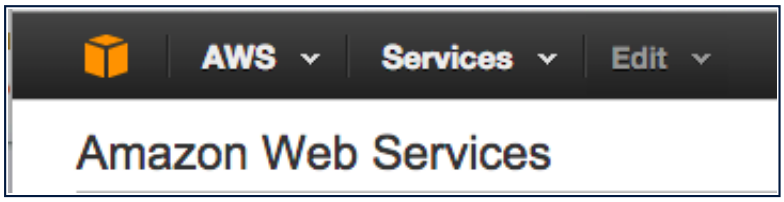

VPC

Under header Networking, select VPC:

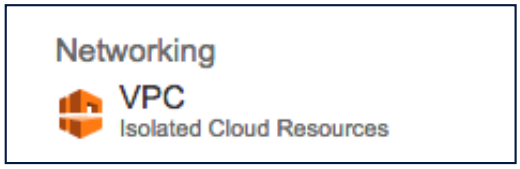

A **virtual private cloud (VPC)** will determine who and what gets to access the site. We will use the wizard and content ourselves with only on VPC. In an enterprise-level application, you will want at least 4, 2 to be private and run your database, and two to be public and hold your web-serving application. By duplicating the private and public VPCs you can benefit from fail-over and load balancing tools. For our purposes, we'll just go with one public VPC:

Start the "wizard":

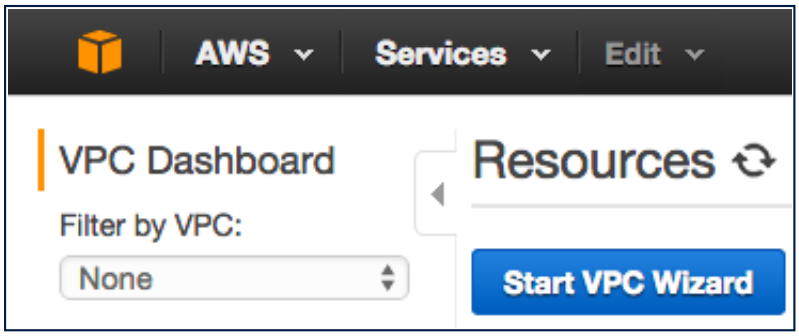

Start the wizard and select "VPC with a Single Public Subnet":

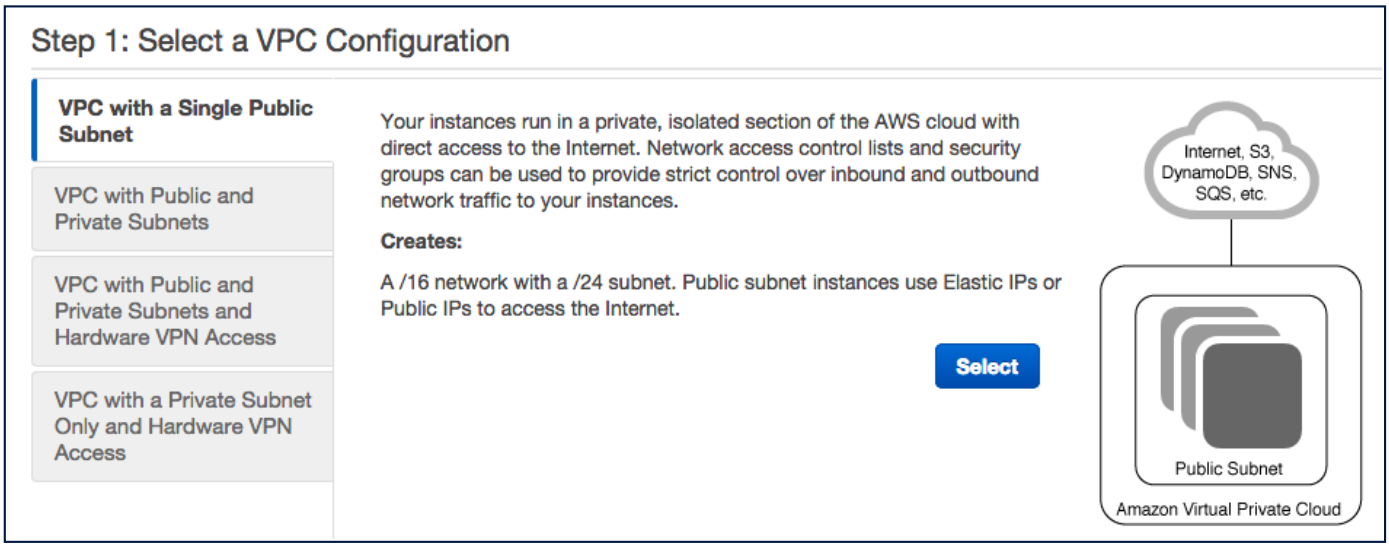

Most of the defaults are fine except you should add a name under “VPC name”:

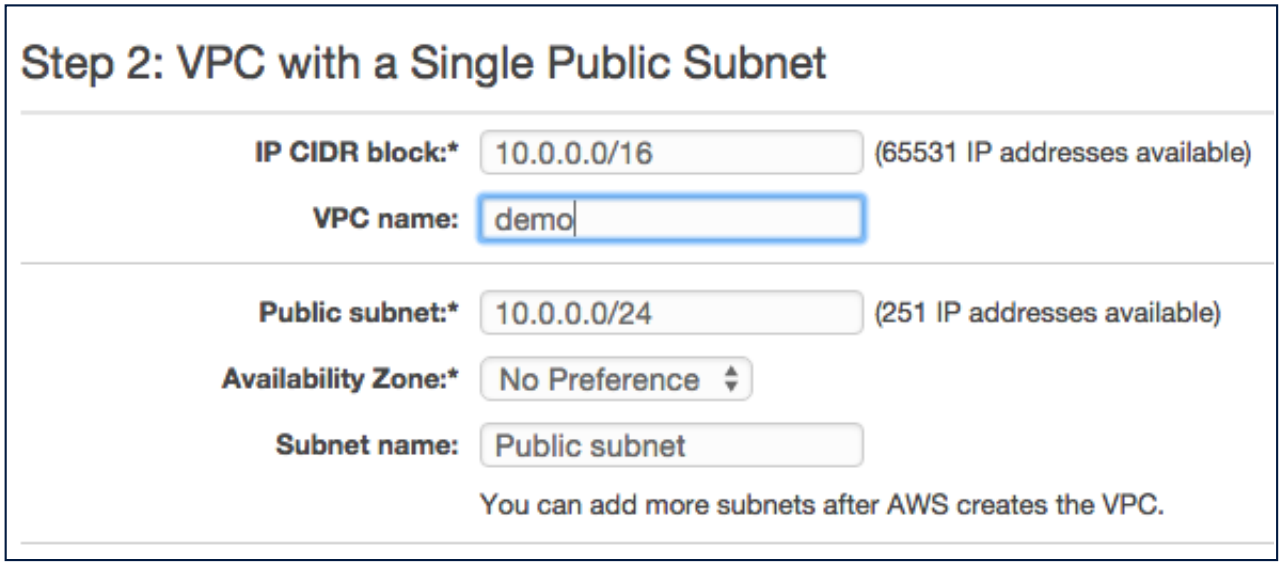

EC2

VPC is done, let’s now create our EC2 instance - this is going to be our cluster-launching machine. Click on the orange cube in the upper left corner of the page. From the ensuing menu, choose the first option, “EC2”:

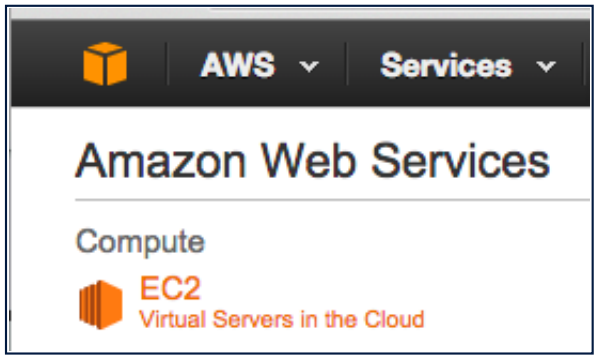

In `Create Instance`, select `Launch Instance`:

Select the first option `Amazon Linux AMI`:

In **Step 2**, continue with the preselected machine and click `Next: Configure Instance Details`:

In **Step 3**, keep all defaults but change the `Auto-assign Public IP` to `Enable`:

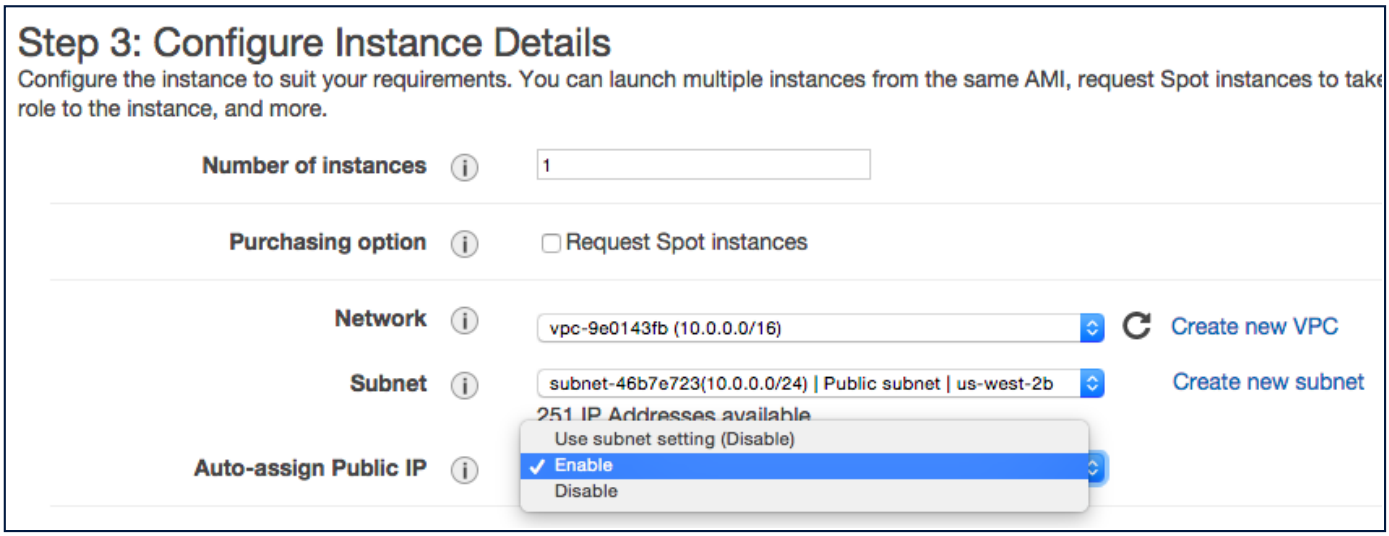

Select `Review and Launch`:

In **Step 7**, `SSH port 22` is opened by default so there is nothing for us to do but to click `Launch`. It will open a pop-up box where you will need to create a new key pair:

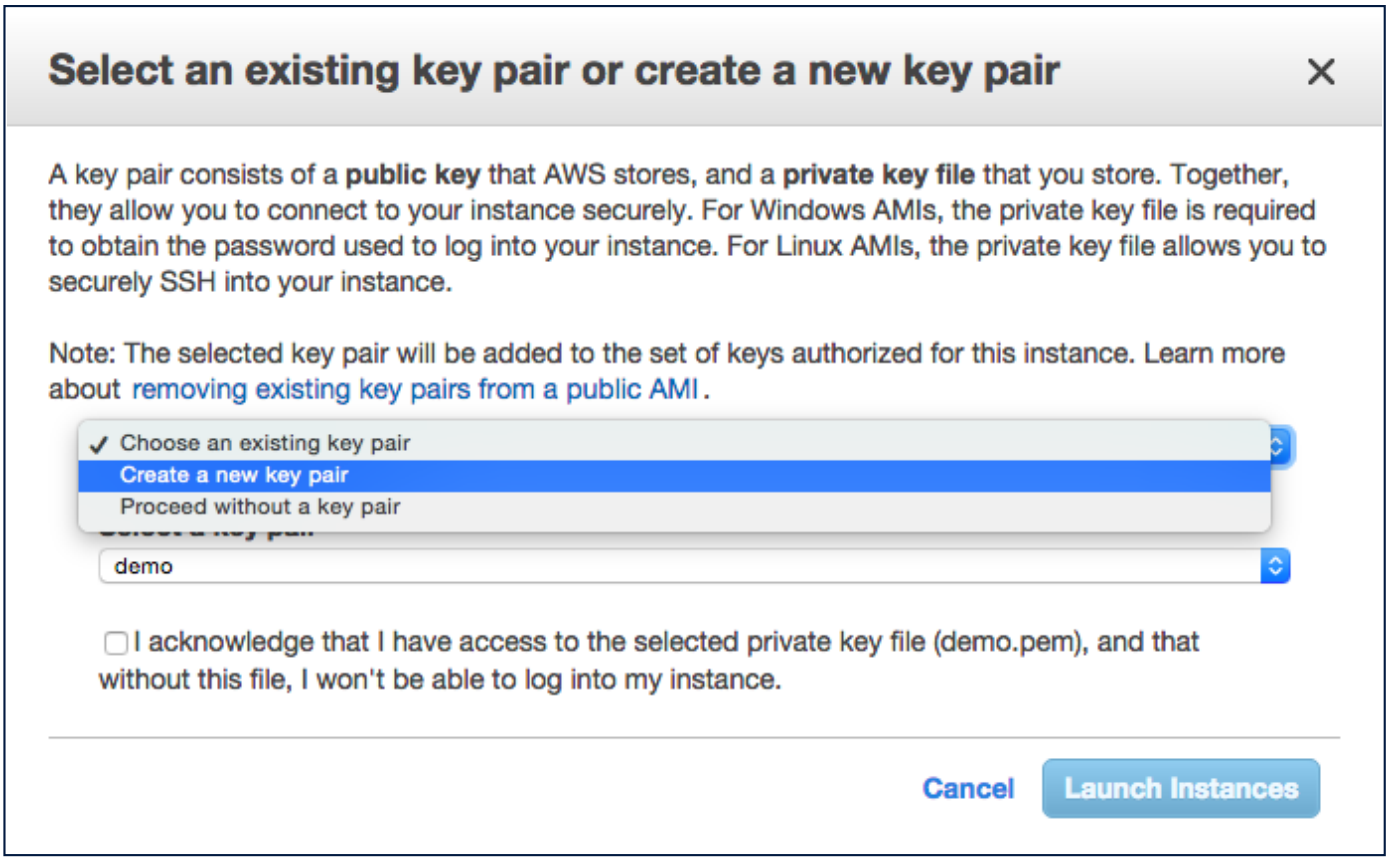

Key-pair is a security file that will live on your machine and is required to SSH into the instance. I tend to create them and leave them in my downloads. What ever you decided to do, make sure you know where it is as you'll need to pass a path to it every time you want to connect to it. Create a new one, check the acknowledgment check-box.

Finally, click Launch Instances. That's it for our instance. Hit the View Instances on the next page as the machine is being initialized:

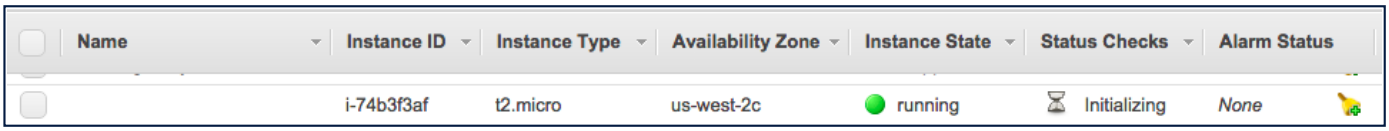

We need to wait till the Initializing is done. Meanwhile we can get the security credentials that will be needed to launch the clusters. Go to your account name drop down in the top right corner and click `Security Credentials`:

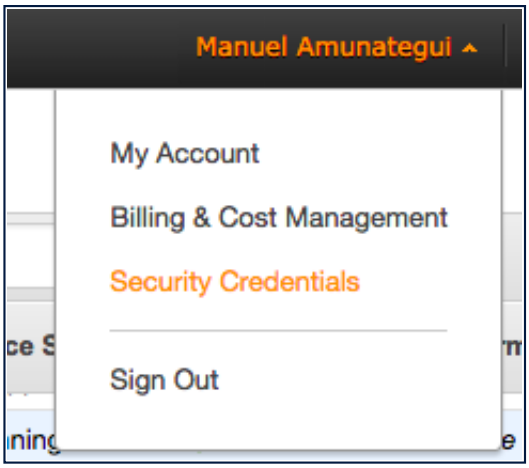

Choose `Access Keys`

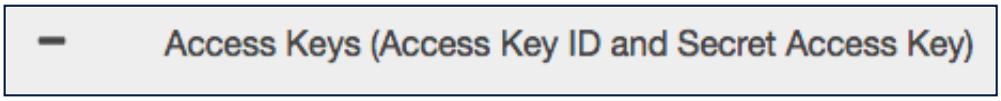

And `Create Access Key`

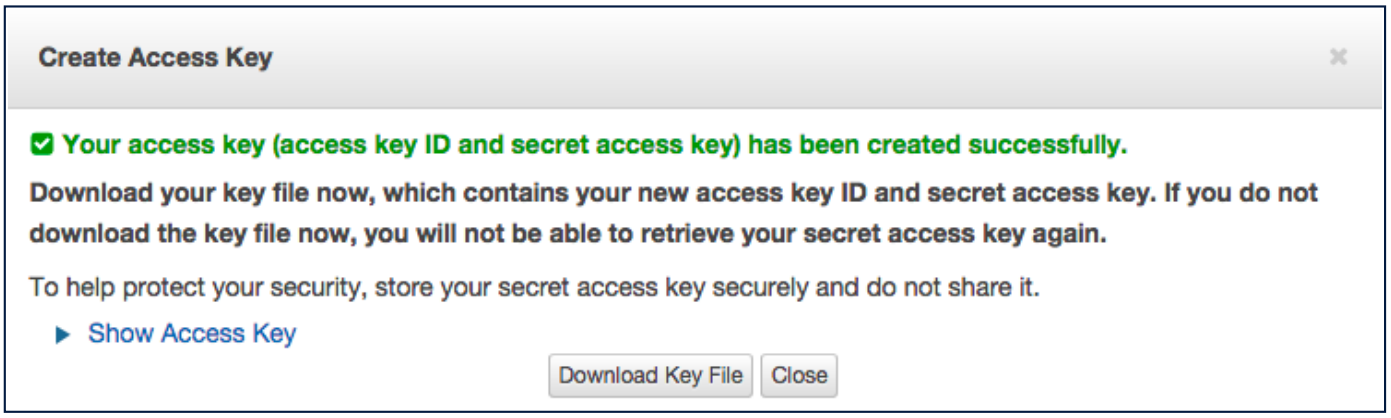

This will download two strings:

```
AWS_ACCESS_KEY_ID=...
AWS_SECRET_ACCESS_KEY= ...
```

Keep that download handy as we'll need the values shortly. If you need more information on security keys check out this Amazon blog post

Back at the instance window, check that our new instance is running, you can click on the check box to access the assigned dynamic IP address:

Connecting to EC2

Use the `Actions` button to get the exact SSH string to reach the instance:

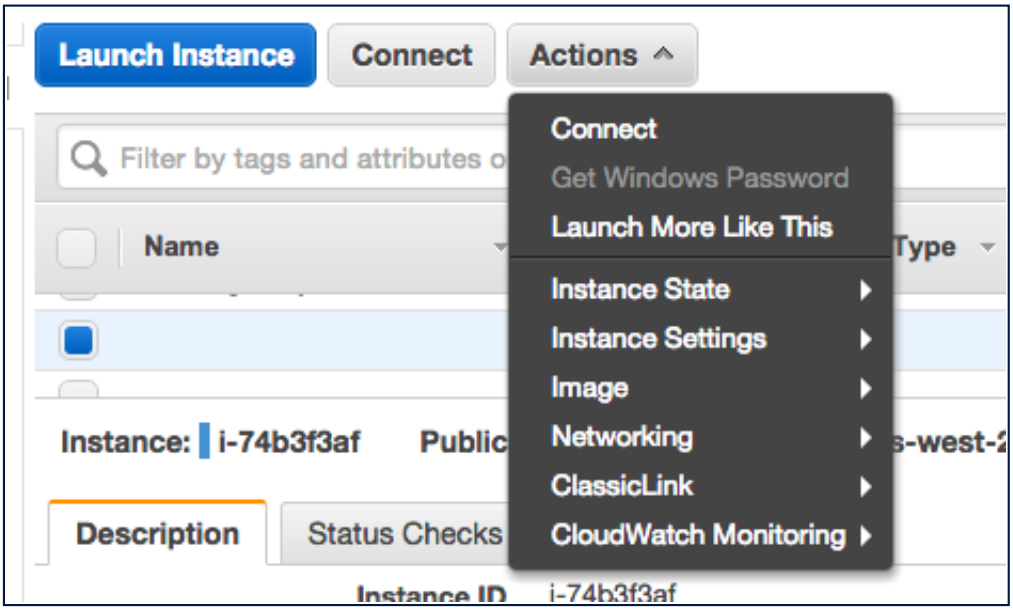

And select `Connect`:

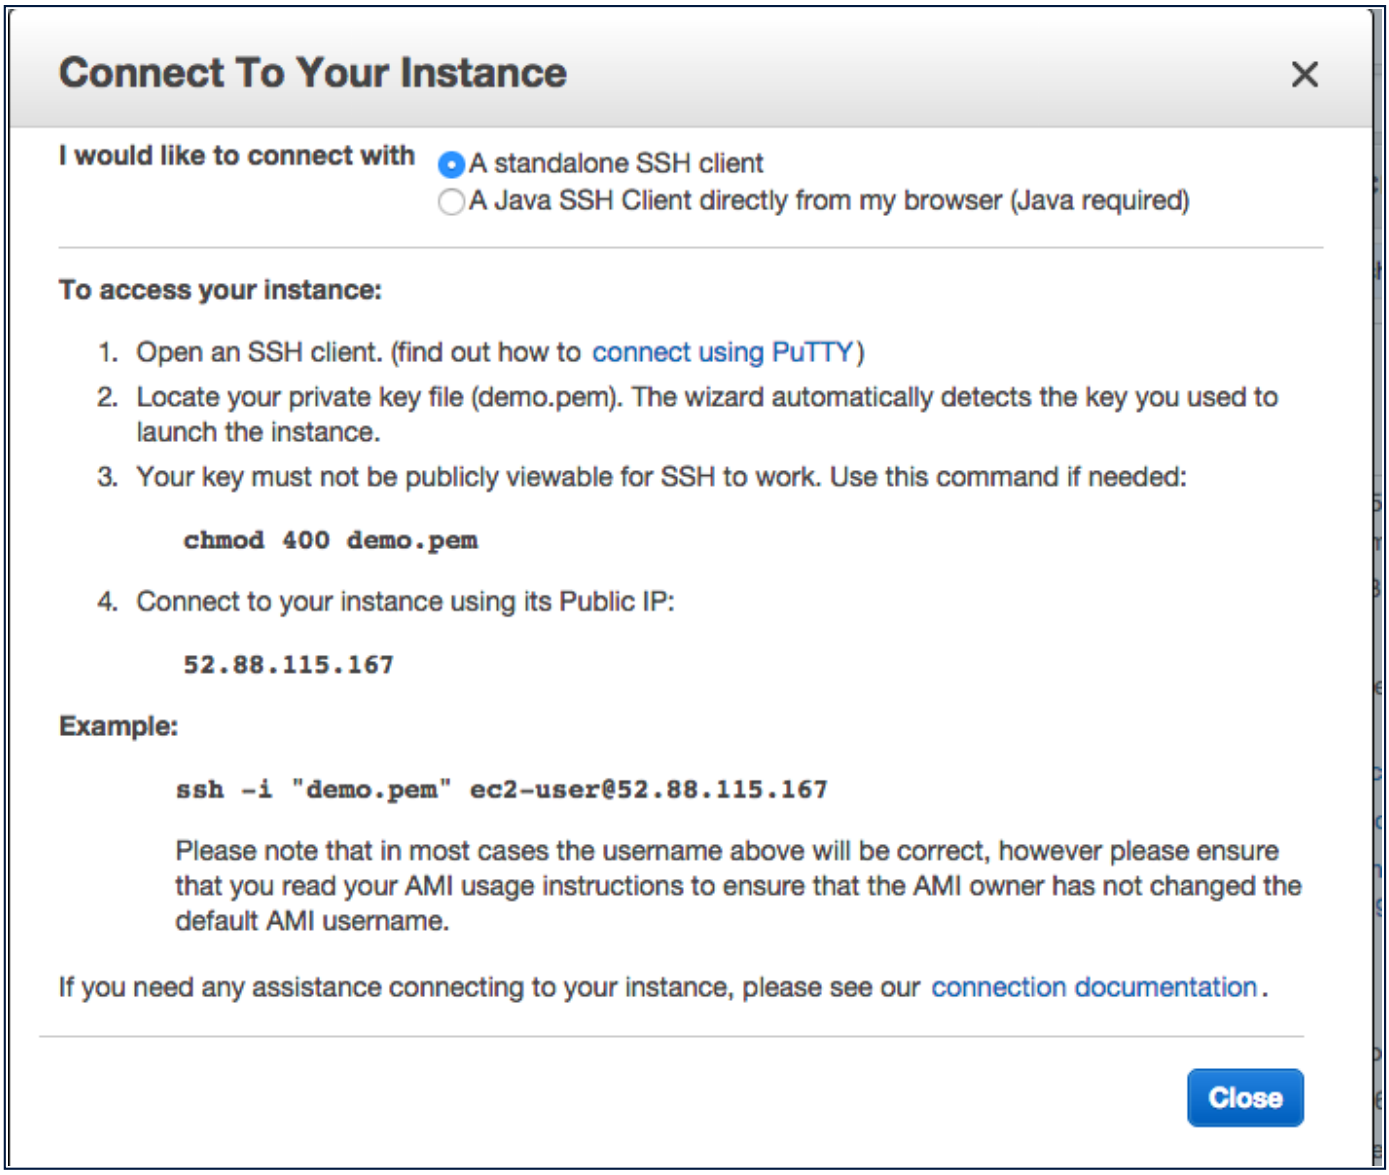

Read the instruction on the pop-up box. The last line states your connection string: `ssh -i "udemy.pem" ec2-user@52.88.115.167`. To use it on the Mac, open your terminal and navigate to your Downloads folder (or wherever you moved your pem key-pair file). As per instructions, apply `chmod 400 udemy.pem` and copy the example string.

SSH

Paste it into the terminal and follow the instructions to log into your EC2 instance:

```
ssh -i "udemy.pem" ec2-user@52.88.115.167
```

Your terminal should confirm the EC2:

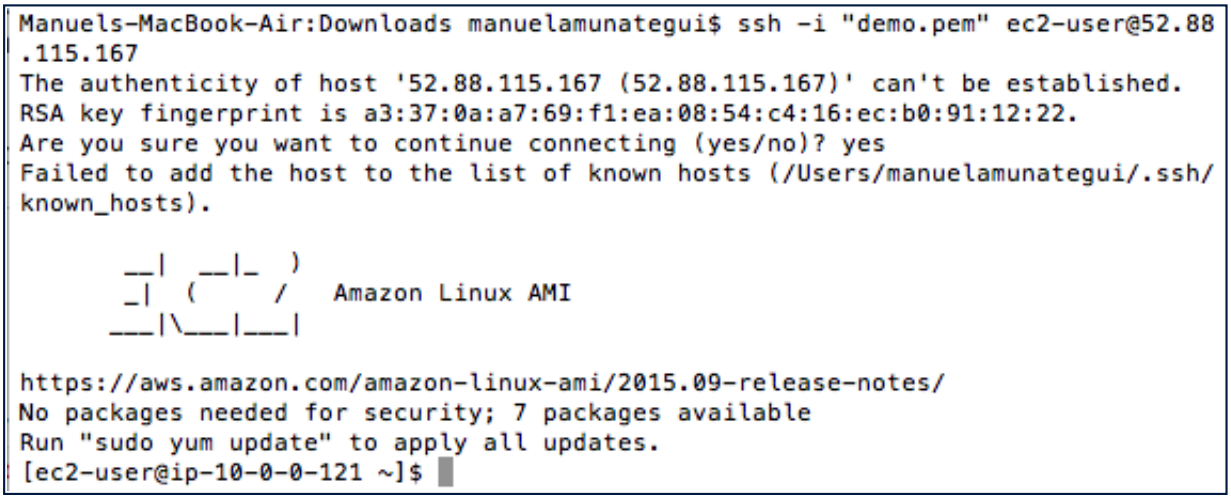

Next we will get **Spark clusters** up and running.